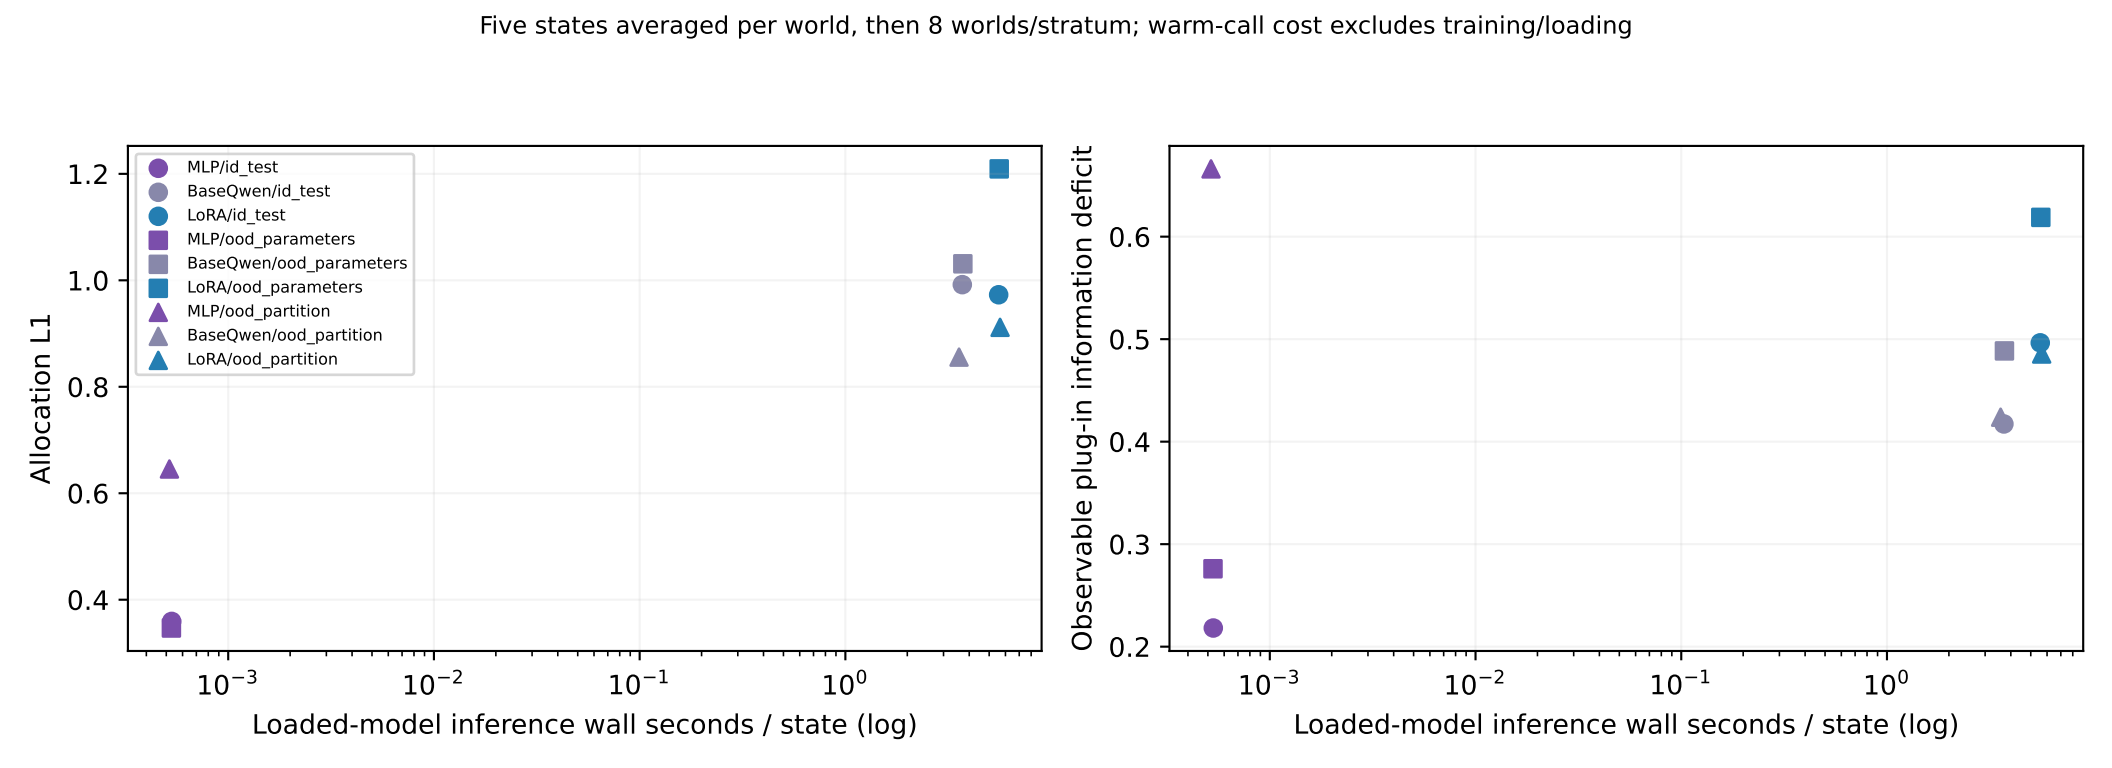
Five states averaged per world, then 8 worlds/stratum; warm-call cost excludes training/loading
MLP/id_test
BaseQwen/id_test
LoRA/id_test
MLP/ood_parameters
BaseQwen/ood_parameters
LoRA/ood_parameters
MLP/ood_partition
BaseQwen/ood_partition
LoRA/ood_partition
Allocation L1
Observable plug-in information deficit
Loaded-model inference wall seconds / state (log)
Loaded-model inference wall seconds / state (log)